

APAS-VERUS: AI Paired Proof Engineering Techniques and Experience

- ▶ Brian G. Milnes briangmilnes@gmail.com
- ▶ Experience, Results and Techniques in building
- ▶ Algorithms Parallel and Sequential by Acar and Blelloch
- ▶ in Rust
- ▶ and then proving it in Verus.
- ▶ Or **What Can We Prove Today and for How Much?**
- ▶ <https://github.com/briangmilnes/{APAS-VERUS,APAS-AI,rusticate,veracity}>

Outline of the talk

- ▶ Its a pleasure to speak here at CMU CS, my intellectual home.
- ▶ Background
- ▶ Algorithms Parallel and Sequential (APAS)
- ▶ Rust - The Good, The Bad and the Ugly
- ▶ APAS-AI - AI Paired Programming APAS in Rust
- ▶ Rusticate - sending Python back to the family estate
- ▶ Verus - Proving Rust

Outline of the talk

- ▶ APAS-VERUS - AI Paired Proving APAS in Verus
- ▶ Veracity - Software Engineering AI Paired Proving
- ▶ Software Engineering in AI Paired Proving
- ▶ CPR\$ - how to measure costs and mine
- ▶ The AI Paired Programming Interfaces
- ▶ The Internet Apocalypse
- ▶ Review of the Talk
- ▶ Questions

Background

- ▶ B.Sc. Applied Math (Computer Science), Carnegie Mellon
- ▶ 2 years of AI PL development (Carnegie Group), my first startup.
- ▶ 7 years of (symbolic) AI Research (Soar Group, CMU CS)
- ▶ 2 years of PL research (Fox Group, CMU CS) aimed at making type safe languages work in systems and networking.
- ▶ Founding member of Lycos - systems, performance, ran operations
- ▶ Early Amazon - systems, performance, operations
- ▶ M.Sc. Computer Science, University of Washington
- ▶ Early Zillow - 6th engineer, ran operations, systems, performance
- ▶ A general love of proving programs with Rocq, F* and now Verus.

Algorithms Parallel and Sequential (APAS)

- ▶ I chose to implement: Algorithms Parallel and Sequential (APAS) in Verus.
- ▶ Umut Acar and Guy Blelloch 2022
 - ▶ 121 algorithms!
 - ▶ 81 of which can be parallel.
 - ▶ 740 concepts.
- ▶ Not only is it a great textbook but it proved with very few changes needed!
- ▶ Thank you Umut and Guy!

Rust - The Good

- ▶ 20 year old PL
- ▶ Fast compilation, code.
- ▶ Industrial acceptance: AWS, Google, Huawei, Microsoft, and Mozilla.
- ▶ Linux Kernel now uses it!
- ▶ Linear typing plus borrowing! No GC!
- ▶ Clear mutability.
- ▶ Slicing! with ownership.
- ▶ Their 'cargo' package system works rather well.
- ▶ No objects for modularity! Good.

Rust - The Bad

- ▶ No GC! Circular structures require GC and why not have it!
- ▶ No regions, just lifetime with end of scope deallocation which is slower.
- ▶ Translating C to Rust is hard to get it into linear logic + borrowing.
- ▶ However, you can (and I have) rewritten algorithms with a free list (that are right).
- ▶ Macros, with typing checked at use, which is not so good.
- ▶ No objects: it's great for what it was invented for: modeling the real world and interfaces by Ivan Sutherland in 1962! (sorry Bob!).
- ▶ Rust terminology is random.

Rust - Typeclasses are a weak module

- ▶ You get basic module with `pub/pub(crate)` and `no-pub`.
- ▶ You get typeclasses which **MUST** be implemented at type.
- ▶ The Rustaceans have decided that unless their module implements a typeclass at multiple types, they won't use it.
- ▶ So reading Rust is every bit as scattered as reading C.
- ▶ Verus is now doing this also, but I **ABUSE** the notation to put all the specs together in APAS-VERUS for readability.
- ▶ You will pine for ML modules!

Rust - The Ugly — Equality and Ordering

Property	PartialEq	Eq	PartialOrd	Ord
Reflexive	NO!	req	NO! on leq	yes
Symmetric	req	req		
Transitive	req	req	req	req
Antisymmetric			req	req
Total				req
Consistent w/ ==				

Rust - The Ugly

- ▶ Clone
 - ▶ `fn clone(&self) -> Self`
 - ▶ Informal contract: returns a value equal to `*self`
 - ▶ No hard language enforcement — you can implement a “clone” that returns something different , but it violates the convention.
- ▶ Copy: Clone
- ▶ Copy is a subtrait of Clone — every Copy type must implement Clone
- ▶ Hard contract: `clone()` must be equivalent to a bitwise copy, i.e., `clone() == *self`
- ▶ This is documented in std: “if `T: Copy`, `T::clone(&x)` must be equivalent to copying `x`”
- ▶ What you really want here is to detangle these.

APAS-AI - AI Paired Programming APAS in Rust

- ▶ APAS-AI is a nearly complete, idiomatic Rust implementation of the algorithms from Acar and Blelloch.
- ▶ Sequential and parallel variants throughout.
- ▶ Timeline: 347 commits over 88 calendar days (Aug–Oct 2025).
- ▶ 59 days of active development; 8 residual commits in November.
- ▶ I knew no AI paired programming when I started.
- ▶ I knew no Rust when I started.
- ▶ My AI had to teach it to me, which was harder than I thought as of the 94 terms used in the Rust language and docs, only 4 are real PL terms!

APAS-AI - AI Paired Programming APAS in Rust

- ▶ Scale:
 - ▶ 42 chapters
 - ▶ 238 source files
 - ▶ 45,485 source LOC.
- ▶ Tests:
 - ▶ 246 files
 - ▶ 55,223 LOC
 - ▶ 3,923 test functions
 - ▶ 1.2× the source code size which is heavy.
 - ▶ But tests are cheap now.
- ▶ Benchmarks: 171 files, 13,890 LOC, 360 benchmark functions

Rusticate - sending Python back to the family estate

- ▶ Built because Python + regex cannot reliably parse Rust.
- ▶ Webster's definition of rusticate:
 - ▶ 1. To go to or live in the country.
 - ▶ 2. (British) To suspend a student from a university, especially Oxford or Cambridge, as a disciplinary punishment.
- ▶ Rusticate is a suite of 89 tools, most deprecated, for analyzing and transforming rust codebases.
- ▶ review tools, fix tools, metrics, and migration aids.
- ▶ The string hacking detector is the foremost valuable tool to fix any and all transformation of Rust.

Verus: Verified Rust

- ▶ Verus is a tool for formally verifying Rust, developed at MSFT Research, VMware Research and Carnegie Mellon University.
- ▶ Team: Andrea Lattuada, Travis Hance, Chris Hawblitzel, Jay Lorch, Matthias Brun, Chanhee Park, Yi Zhou, Jon Howell, Bryan Parno.
- ▶ Goals: bring machine-checked proof to systems software written in Rust.
- ▶ Goals: tight integration with the programming language.
- ▶ Goals: easier faster proof.

Key Verus Publications

- ▶ “Verus: Verifying Rust Programs using Linear Ghost Types”
Lattuada, Hance, Cho, Brun, Subasinghe, Zhou, Howell, Parno, Hawblitzel
- ▶ “Verus: A Practical Foundation for Systems Verification”
Lattuada, Hance, Bosamiya, Brun, Cho, LeBlanc, Srinivasan, Achermann, Chajed, Hawblitzel, Howell, Lorch, Padon, Parno
- ▶ “Verifying Concurrent Systems Code” Travis Hance — PhD Thesis, Carnegie Mellon University, 2024

Verus

- ▶ Specifications are written in a pure math sublang: First order logic.
- ▶ spec functions, forall/exists, arithmetic, sets, sequences.
- ▶ Undecidable.
- ▶ Z3 handles linear arithmetic, arrays, and quantified formulas, but quantifier instantiation requires explicit trigger annotations.
- ▶ But there is also a faster linear arithmetic solver, Singular.
- ▶ Wraps existing Rust code:
 - ▶ spec / proof / requires / ensures on fns
- ▶ Ships the Rust binaries directly.
- ▶ Linear Logic + Borrowing from the Rust type system, which rustc checks.
- ▶ The TCB then is Verus+Z3+RustC+STD, which is BIG.

Views and the Libraries

- ▶ A View maps an executable Rust type to a mathematical ghost type: “Vec views as Seq”, “HashSet views as Set”.
- ▶ Specs are written over the view; exec code manipulates the real type.
- ▶ vstd is Verus’s standard library — specs for Vec, Seq, Set, Map, Multiset, arithmetic, common lemmas, Fns ...
- ▶ Ghost types live only in the verifier — they have no runtime cost and no runtime representation.

UnDirGraphStEph — Data Struct + View

```
#[verifier::reject_recursive_types(V)]
pub struct UnDirGraphStEph<V: StT + Hash> {
    pub V: SetStEph<V>,
    pub E: SetStEph<Edge<V>>,
}

impl<V: StT + Hash> View for UnDirGraphStEph<V> {
    type V = GraphView<<V as View>::V>;
    open spec fn view(&self) -> Self::V {
        GraphView { V: self.V@, A: self.E@ }
    }
}
```

UnDirGraphStEph — Trait

```
pub trait UnDirGraphStEphTrait<V: StT + Hash>:  
  View<V = GraphView<<V as View>::V>> + Sized  
{  
  open spec fn spec_neighborhood(&self, v: V::V) -> Set<V::V>  
    recommends spec_graphview_wf(self@), self@.V.contains(v)  
  { Set::new(|w| self@.A.contains((v,w)) || self@.A.contains((w,v))) }  
  
  fn neighborhood(&self, v: &V) -> (nbrs: SetStEph<V>)  
    requires spec_graphview_wf(self@), self@.V.contains(v@), ...  
    ensures  nbrs@ == self.spec_neighborhood(v@), nbrs@ <= self@.V;  
}
```

UnDirGraphStEph — Impl: neighborhood exec body

```
fn neighborhood(&self, v: &V) -> (nbrs: SetStEph<V>) {  
  let mut nbrs: SetStEph<V> = SetStEph::empty();  
  let mut it = self.E.iter();  
  let ghost edges_seq = it@.1;  
  loop  
    invariant nbrs@ == Set::new(|w| exists |i|  
      0 <= i < it@.0 && /* edges_seq[i] hits v on either side */ ...),  
    decreases edges_seq.len() - it@.0,  
  {  
    match it.next() {  
      None    => { proof { /* set-equality lemmas */ } return nbrs; }  
      Some(e) => { /* if feq(ℓe.0,v) insert e.1; sym. */ ... }  
    }  
  }  
}
```

Wrapping Rust — Declaring an external type

- ▶ Four specification constructs are used to give specs to Rust stdlib.
- ▶ A proxy struct that introduces a spec for a foreign type.
- ▶ The proxy struct name is conventionally ExTypeName.

```
#[verifier::external_type_specification]  
pub struct ExVec<T>(Vec<T>);
```

Wrapping Rust — Declaring an external function/method

- ▶ A proxy function with the same signature as the foreign function, carrying the requires/ensures.

```
#[verifier::external_fn_specification]
pub fn ex_vec_push<T>(v: &mut Vec<T>, value: T)
    requires v@.len() < usize::MAX,
    ensures  v@ == old(v)@.push(value),
{ v.push(value) }
```

Wrapping Rust — Declaring an external function/method

- ▶ Add a View and specs to a foreign type.

```
#[verifier::external_type_specification]
pub struct ExHashMap<K, V>(HashMap<K, V>);
impl<K,V> View for ExHashMap<K,V> {
    type V = Map<K::V, V::V>;
    spec fn view(&self) -> Map<K::V, V::V>;}
```

Wrapping Rust — external_trait_specification

- ▶ Adds a spec to a foreign trait without modifying it:

```
#[verifier::external_trait_specification]
pub trait ExClone: Sized {
    type ExternalTraitSpecificationFor: core::clone::Clone;
    fn clone(&self) -> Self;
}
```

- ▶ The proxy trait name is Ex<TraitName> by convention.
- ▶ Add requires/ensures to the method to give it a full contract.
- ▶ Limitation: no generics.

Tokenized State Machines — Hance, CMU 2024

- ▶ Problem: Rust's ownership types handle sequential aliasing well but cannot express distributed protocol state across threads.
- ▶ Answer: A Tokenized State Machine defines protocol state as fields with sharding strategies (variable, map, count, storage_option...).
- ▶ Transitions and an inductive invariant are proved once, globally.
- ▶ Verus auto-generates ghost token types and exchange functions

Tokenized State Machines — Hance, CMU 2024

- ▶ So client code manipulates local tokens, not global state.
- ▶ RwLock in vstd is implemented via a tokenized state machine
- ▶ The lock's internal protocol (unlocked / read-locked / write-locked, reader count) is the TSM.
- ▶ RwLockPredicate is the (single) invariant the user supplies.

Tokenized State Machines — Hance, CMU 2024

- ▶ Hance et al. “Sharding the State Machine” — OSDI 2023 (primary TSM paper)
- ▶ Hance, Howell, Padon, Parno. “Leaf” — OOPSLA 2023 (storage protocols)
- ▶ Hance. PhD Thesis, CMU-CS-CS-24-146, 2024 (full formal treatment)

Verus: How Fast Verus Is Moving

- ▶ 4,225 commits since March 2021 — 5 years, still accelerating
- ▶ Commits per year:
- ▶ 2021: 474 2022: 906 2023: 992
- ▶ 2024: 745 2025: 761 2026: 347 (thru April)
- ▶ Rolling releases: 396 tagged — roughly one every 4 days
- ▶ 59 stable releases!
- ▶ Excellent team and work. I have had only a few crashes, which I could work around.

APAS-VERUS - AI Paired Proving APAS in Verus

- ▶ Goal: formally verify all algorithms in Acar and Blelloch
 - ▶ Every algorithm gets a machine-checked proof
 - ▶ no admitted lemmas,
 - ▶ no hand-waving in production code.
- ▶ 44 chapters, 262 algorithm files, upto 4 variants per algorithm: StEph (sequential mutable), StPer (sequential persistent), MtEph (parallel mutable), MtPer (parallel persistent).
- ▶ Minimal use of Rust std.
- ▶ But as you'll see I had to write some axioms and admit a class of functions.

APAS-VERUS - AI Paired Proving APAS in Verus

- ▶ 26 vstdplus library modules, mostly making Views.
- ▶ 29 standards documents encoding project proof conventions,
- ▶ Verification is the primary goal.
- ▶ Runtime tests (RTT) and proof-time tests (PTT) are secondary, but still terribly useful.
- ▶ They both still catch errors and instruct the AI.

APAS-VERUS — Quantitatives

- ▶ Scale: 44 chapters, 262 files, 186,223 src LOC (not counting comments).
- ▶ With vstdplus, standards, RTT, PTT: 275,014 total LOC.
- ▶ Built in 160 days, 2,596 commits, 8 agents, 281 agent-round reports.
- ▶ Verification: 5,674 verified proof obligations, 0 errors.

APAS-VERUS: Full Validation Cost (2026-04-12)

- ▶ Elapsed: 210s
- ▶ rust_verify RSS: 10,278 MB (~10 GB)
- ▶ Z3 RSS: 6,874 MB (~6.7 GB)
- ▶ rust_verify CPU: 216s
- ▶ Z3 CPU: 265s
- ▶ But I have had Z3 jump up to as much as 28 GB when I write bad proofs.
- ▶ Somewhere in it I suspect there is a novel formal verification of some algorithm.

APAS-VERUS — Quantitatives

- ▶ Runtime tests: 3,776 pass in 21 s.
- ▶ Verus has a nice proof-time test harness, so I pulled it out:
 - ▶ 221 pass in 259 s.
 - ▶ I used it mostly to continuously track that my iterators prove and continue to prove with verus changes and my specifications.
 - ▶ These caught dozens of problems that would have been conflated with algorithms loops.
- ▶ Benchmarks in 42s.

APAS-VERUS — Quantitatives

- ▶ Holes: started at 238 (R20), now 0!
- ▶ Largest chapter is the forest: Chap37 - AVL trees, BST variants - 20,319 src LOC.
- ▶ $2 \times$ more source code to verify than APAS-AI needed to implement.
- ▶ Start: 2025-11-03
- ▶ End : 2026-04-12
- ▶ Duration: 150 calendar days, roughly (more later)

APAS-VERUS -

- ▶ Spec 32,868 (21%)
- ▶ Proof 42,251 (27%)
- ▶ Exec 67,883 (44%)
- ▶ Rust 12,206 (8%) plain Rust (outside Verus!)
- ▶ The “rust” 8% is code outside Verus! — Debug impls, macros, cfg, etc.
- ▶ All proof to exec: $75,119 / 67,883 = 110\%$.

APAS-VERUS: The Pain Points

- ▶ When I started Verus, iterators for collections took quite some time.
- ▶ Generics and Equality was the second big pain point.
- ▶ I still have full equality axioms for generic types.
- ▶ Ordering was and is still difficult, it made my UnionFind consume up to 28GB in Z3.
- ▶ Verus is so fast, even with bloated AI proofs, I didn't profile much until I was in the last few chapters.
- ▶ I simply made validate isolate by chapters.

APAS-VERUS: The Pain Points

- ▶ Closures took a good bit of work.
- ▶ Applying them in map/reduce and so on took a while.
- ▶ Rust has $\text{FnOnce} < \text{FnMut} < \text{Fn}$, but no real FPure .
- ▶ You can `const fn` and the compiler will check it but then you can't say: `fn map(const fn F)`
- ▶ Verus Ghost functions allows good validation but purity would be simpler.

AutoCLRS

- ▶ AutoCLRS is Swamy et al. and Al's implementation: "Introduction to Algorithms, 4th ed" 2022
- ▶ by Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest, and Clifford Stein.
- ▶ in Pulse.
- ▶ RISE MSR blog (2026-03-06) says the initial 10K lines came "very quickly" and then "about a month of nudging" to reach 100K LOC.
- ▶ And now seems to be about 130K LOC.
- ▶ Nikhil Swamy with thanks to Gabriel Ebner, Lef Ioannidis, Guido Martinez, Matthai Philipose and Tahina Ramananandro.

AutoCLRS

- ▶ They definitely had some tool advantages in terms of incremental proofs through a server.
- ▶ That's another verus pain point but not too bad.
- ▶ Plus, they knew F* and Pulse to start!
- ▶ And they did a formal specification of algorithmic complexity!

Veracity- Software Engineering AI Paired Proving

- ▶ Veracity is a suite of 22+ tools for analyzing, reviewing, and fixing Verus codebases.
- ▶ Review tools:
 - ▶ proof holes (assume, external_body, admit),
 - ▶ style enforcement (21 rules, auto-reorder),
 - ▶ with spec strength classification fed to AI,
 - ▶ veracity-count-loc (spec/proof/exec breakdown),
 - ▶ chapter-cleanliness-status (clean vs. holed chapter summary vs blocked by),
 - ▶ string-hacking detector, function inventory, etc.
- ▶ One of the best is veracity-minimize-proofs

Veracity- Software Engineering AI Paired Proving

- ▶ All tools are AST-aware (`ra_ap_syntax` / `Verus_syn`).
- ▶ A string-hacking detector flags usage of string manipulation instead of AST work.
- ▶ And when bugs appeared they were mostly string hacking.
- ▶ Because no matter what I said to my AIs they LOVE string hacking shortcuts.
- ▶ Heck, I had to run the string hacking detector on the string hacking detector.

Veracity- Software Engineering AI Paired Proving

- ▶ Search: veracity-search — type directed search over vstd
- ▶ VERUS by type signature, finding lemmas before writing new ones.
- ▶ “Specifications as Search Keys for Software Libraries” Eugene J. Rollins and Jeannette M. Wing
- ▶ Written for my sins of asking why does vstd not have X, when it did!
- ▶ Even more useful for my AI’s seriously disturbing sinning.
- ▶ This allowed me to download ALL known Verus (git VerusCodebases) and have my AI search them in 1.2 seconds!
- ▶ Does F*/Pulse have one yet?
- ▶ It only took about a day.

APAS-VERUS - Complexity

- ▶ APAS states complexity and informally proves many of them for some algorithms.
- ▶ I had to build a tool to get the right ones in the code at the right place.
- ▶ My single threaded implementations often don't match the textbook intentionally.
- ▶ Then I wrote a programmatic tool to find and list mismatches.
- ▶ Then I had Claude Opus do it's analysis and compare every function with the textbook's.
- ▶ This found 16 faults in parallel algorithms.

APAS-VERUS: Verified Iteration - Pain Point

- ▶ Iteration in Rust is rather complex.
- ▶ 70 functions on iterator but only 7 functions cover the 90% case.
- ▶ And in Verus it's a bit more complex and takes some time to learn.
- ▶ 10 components required per collection (all inside Verus!)
- ▶ 6 verified loop patterns per collection
- ▶ {loop, for} X {borrow iter, borrow into, consume}
- ▶ Iterators requires one assume!

APAS-VERUS: Verified Iteration

- ▶ Verus forbids adding requires on external trait impls (`std::iter::Iterator`)
- ▶ Hand-rolled iterators need `assume(iter_invariant(self))` in `next()`
- ▶ Everything but that one `assume` is fully proved
- ▶ 44 collections implemented; all carry verified iterators
- ▶ Verus has proof time tests inside, I freed them to run in APAS-VERUS.
- ▶ This was critical to get iterative loops to prove over my collections ADTs.
- ▶ You can prove full iterators in your own copy of the traits with no `assume` and I have.

APAS-VERUS: Experiments

- ▶ Agents often say “Verus Can’t Do That”
- ▶ I said “Make an experiment!”
- ▶ Quantitatives:
 - ▶ 168 experiment files
 - ▶ 21,476 lines of code (not counted in the totals)
- ▶ Topics span: Clone, Arc, RwLock, TSM, closures, iterators, generics, float, bitvector, PartialEq, Copy, async, hash tables, parallel algorithms, ghost types, Send/Sync, collect, and sorting.

APAS-VERUS: Experiments

- ▶ Results:
- ▶ 107 files: SUCCEEDS / VERIFIES — pattern adopted into codebase
- ▶ 61 files: FAILS — Verus limitation documented, workaround noted
- ▶ Notable successes that unlocked chapters:
 - ▶ TSM/RwLock layer pattern — unlocked all Mt modules,
 - ▶ Named closure ensures through ParaPair — unlocked fork-join,
 - ▶ Ghost struct Send/Sync — unlocked AVLTreeSetsMtEph,
 - ▶ Tree module style — unlocked Kruskal, Prim, UnionFind.

APAS-VERUS Standards

- ▶ I finally built a set of coding standards in Verus rust files and in comments.
- ▶ Agents read all standards before every task (~6,200 lines, ~54K tokens)
- ▶ Violations are mostly AI checked except where an extensive code styling can get things.
- ▶ Quantitatives:
 - ▶ 29 standard files
 - ▶ 6,911 lines total
- ▶ Doing this earlier would have really sped things up.
- ▶ My CLAUDE.md would just not do enough even with 50KB and 13K tokens.

APAS-VERUS RULEs

- ▶ Question what are your favorite AI rules?
- ▶ Mine are:
 - ▶ Don't Over Think,
 - ▶ DISCUSS,
 - ▶ Don't jump ahead,
 - ▶ go step by step,
 - ▶ PBOGH: Prove Big or Go HOME!
- ▶ You just have to tell the agents keep on proving!
- ▶ Are you building big CLAUDE.md?
- ▶ Or Cursor Rules?

Veracity: Als Write Redundant Proofs

- ▶ AI proof agents produce many correct but bloated proofs
 - ▶ redundant asserts, unnecessary proof blocks
- ▶ They verify, but they waste solver budget on every subsequent run.
- ▶ So I wrote a proof minimizer: veracity-minimize-proofs.
- ▶ It tests each assert and proof block individually: removes it, re-verify, comment it out if it is not needed and it does not increase time or memory.
- ▶ Result across APAS-VERUS: 22 asserts and 33 proof blocks removed in 105 minutes of wall time. 55 redundant proof statements eliminated, ~2 minutes of minimizer time per removal.


Veracity: Als Write Redundant Proofs

- ▶ One assert in Chap43 OrderedTableMtEph saved:
 - ▶ 104 s of Z3 CPU
 - ▶ up to 89 MB of Z3 RSS per verify run.
- ▶ Eight removals in that one file: Z3 RSS dropped by 57%.
- ▶ 105 minutes of running the minimizer bought many hours of validation drop.
- ▶ This might be novel. Anyone know of anything except Isabelle's sledgehammer?

Veracity Annotations


- ▶ I ended up having to have my tools work mostly in comments.
- ▶ I added `accept(P)` to mark assumes I allowed, almost all `eq/partialEq/clone`.
- ▶ This can be simplified with some Verus language syntax, but then the AIs are not trained on the symbols.
- ▶ 6,081 // Veracity: NEEDED assert
- ▶ 4,681 // Veracity: NEEDED proof block
- ▶ 1,502 // Veracity: NEEDED assert (speed)
- ▶ 245 // Veracity: NEEDED proof block (speed)—
- ▶ Total annotations: 14,140

Proof Holes Over Time — R20 to R201




Proof Time and Memory — Key Reductions

Key Proof Cost Reductions



Full Codebase Z3 Memory Over Time



Proof Time and Memory — Key Reductions

- ▶ Five techniques were used to optimize
 - ▶ minimize-proofs
 - ▶ profiling
 - ▶ splitting specs and applying them just where they are needed
 - ▶ opaque - to hide a definition within the module
 - ▶ private spec - to hide a definition across modules
- ▶ OrderedTableMtEph: -57% proof time after minimize-proofs (R176).
- ▶ UnionFindPCStEph: 139K Z3 instantiations → 0 after opaque pattern used in a required module.
- ▶ Johnson Chap59: 756 MB → 520 MB (-31%) memory reduction.

CPR\$ — Three Numbers and One Ratio

#	Sym- bol	Name	Measures
1	C	Cost of Code	\$ to produce executable code
2	P	Cost of Proof	\$ to produce specs, contracts, and proofs
3	C+P	Total	Full bill for the verified artifact
4	R	Review Ratio	Fraction of deliverable a proofgrammer must read (LOC0R excluded)

Inputs & Computation

- ▶ $\text{total_hours} = \text{person_days} \times \text{avg_hrs_per_day}$
- ▶ $\text{programmer_cost} = \text{total_hours} \times (\text{programmer_costs} / 1,760)$
- ▶ $\text{AI_cost} = \text{total_hours} \times (\text{ai_costs} / 1,760)$
- ▶ $\text{LoEC_share} = \text{LoEC} / \text{total_LOC}$
- ▶ $C = \text{total_cost} \times \text{LoEC_share}$
- ▶ $P = \text{total_cost} \times (1 - \text{LoEC_share})$
- ▶ $R = \text{LOPC2R} / (\text{LOPC2R} + \text{LOC0R})$

Worked Example — Inputs

- ▶ **Programmer rate:** \$375,000/yr (senior + 50% loading)
 - ▶ **Real AI spend:** < \$7,000 across the whole effort
 - ▶ **Derived --ai-costs:** \$4.886/hr × 1,760 = **\$8,599/yr**
 - ▶ **AI split by task-hours:** 21.6% APAS-AI (\$1,512) / 78.4% APAS-VERUS (\$5,488) | # | Project | Hours | LOC | LoEC | LOPC2R | LOC0R |
- | | |
|--|---|
| :- :————- ————-: ————-: ————-: ————-: ————-: | 1 |
| APAS-AI + Rusticate 309.6 31,751 31,751 — — | 2 |
| APAS-VERUS + Veracity 1,123.4 166,401 56,549 47,828 | |
| 95,908 3 Combined 1,432.96 198,152 88,300 47,828 | |
| 95,908 | |

CPR\$ — Combined Result

#	Quantity	Value
1	C — Cost of Code	\$150,723
2	P — Cost of Proof	\$161,594
3	C + P — Total	\$312,317
4	R — Review Ratio	33.3%
5	\$ / verified deliverable line	\$1.877
6	C / KLOEC	\$2.665
7	P / KLOPC2R	\$3.378
8	P / KLOP	\$6.305

Head-to-Head — seL4 (2009) vs APAS-VERUS

#	Quantity	seL4 (pre-AI)	APAS-VERUS (AI-paired)	Ratio
1	Person-years	22	0.64	seL4 ~34×
2	Hours	45,760	1,123	seL4 ~41×
3	KLOE	10	~57	Verus ~5.7×
4	KLOP	480	~110	seL4 ~4.4×
5	KLOP / KLOE	48	~1.9	seL4 ~ 25 ×
6	Klines / hour	0.0107	0.148	Verus ~14×
7	\$ / KLOE	\$825,000	~\$4,300	seL4 ~ 192 ×
8	\$ / KLOP	\$17,188	~\$2,225	seL4 ~7.7×
9	C + P total	\$8,250,000	\$244,840	seL4 ~ 33.7 ×

Rusticate + Veracity allow quantitative software engineering

- ▶ I downloaded the 1036 most downloaded Rust projects, 3636 crates.
- ▶ I threw out three that wrapped std for asynchrony.
- ▶ I did a greedy set cover analysis of what Rust uses.
- ▶ And another of what verus wraps.
- ▶ This is a golden age for quantitative software engineering.
- ▶ It took just two days of person time.
- ▶ And I never read a line of code in the tools. Just check the output and have the AI write a lot of tests.

Rusticate + Veracity : What Rust Cargos use in std.

- ▶ Top 1000 projects, 3636 crates.
- ▶ 19 data types fully support 90%.
- ▶ 48 data types fully support 100%.
- ▶ 69 modules fully support 95%.
- ▶ 79 modules fully support 99%.
- ▶ 1121 methods fully support 95%.
- ▶ 1733 methods fully support 99%.
- ▶ 14,317 total fn definitions in std/core/alloc in Rust
- ▶ 4,965 pub fn definitions in std/core/alloc in Rust
- ▶ But how many of those private functions do we need? I don't know.

Compare Rust STD to APAS-VERUS?

- ▶ APAS-VERUS: -6,401 exec functions total, -4,911 with proofs
- ▶ But APAS-VERUS has $\{Mt, St\} \times \{Per, Eph\}$
- ▶ So it is much more like 2000 distinct functions.
- ▶ And I wrote this in 160 days while learning Verus.
- ▶ At least 30 of those days were understanding and working around pain points.

Rusticate + Veracity: What Verus Wraps

- ▶ As of 2026 14 April
- ▶ How many Rust Data Types does Verus wrap?
 - ▶ 29 types currently wrapped.
- ▶ How many Rust Traits does Verus wrap?
 - ▶ 147 traits in vstd.
- ▶ How many total Rust Methods does Verus wrap?
 - ▶ 154 methods currently wrapped.
- ▶ A small but growing percentage.

Software Engineering in AI Paired Proving

- ▶ Software engineering in the AI age is much like managing programmers.
- ▶ You can't possibly read all their code.
- ▶ So you use:
 - ▶ linters both programmatic and AI
 - ▶ stylers both programmatic and AI
 - ▶ AI reviews directed by programmatic tools.

Software Engineering in AI Paired Proving

- ▶ Starting with a textbook that has a ton of good prose is a huge plus.
- ▶ I did discover a few things about the textbook.
- ▶ It needs more on Scheduling. I only built Help-First scheduling.
- ▶ It used (K,V) in tree sets as mappings, but I had to make this explicit. Particularly with an Ordered Key Mapping.
- ▶ With this Ordered Tables, which are not too complex, were much easier to prove.
- ▶ APAS's discussion of Union Find was too thin. Path compression was very difficult.
- ▶ I had to start with Nikhil Swammy et. al's AlgoCLRS's implementation and proofs.

Software Engineering in AI Paired Proving

- ▶ I use git work trees for proof.
- ▶ The good news is git is powerful and the agents understand it.
- ▶ The bad news is the agents can indeed start using advanced features and get things wrong.
- ▶ I work with between 1 and 9 agents.
- ▶ One on veracity.
- ▶ One orchestrator, writing plans, prompts and controlling merges.
- ▶ 1-8 on branches each working on different independent plans.
- ▶ I am not nearly comfortable enough with agent proof to let them do subagents without observing.
- ▶ Cursor was much better at interruption to guide things.
- ▶ Claude is getting better at it.

Software Engineering in AI Paired Proving- Models

- ▶ The agents I used are: various Google, OpenAI and Anthropic.
- ▶ A google agent once took 12 minutes to 'echo HI' to test it's terminal.
- ▶ Anthropics models are significantly better than OpenAI.
- ▶ And Claude Code's cost is the least so far.
- ▶ They are said in the press to be spending \$5000 per month per user.
- ▶ They are buying market share (what happened to Amazon's buying market share phase?).
- ▶ And they are buying coding interactions.
- ▶ But both vendors are playing leap frog.

Software Engineering in AI Paired Proving- Problems

- ▶ They are inconsistent.
- ▶ They lie.
- ▶ They cheat.
- ▶ They say I can't prove that!
- ▶ They err.
- ▶ But mostly they are just forgetful.
- ▶ My attempts to give them coding standard checklists, ala Watts Humphrey and the PSP failed.
- ▶ Sometimes they would say "The checklist was followed" and then admit "I lied."

Software Engineering in AI Paired Proving- Problems

- ▶ They are getting better, very few code hallucinations.
- ▶ If you tell them you are Dr. FunkenProof and they are Igor, they hallucinate whole modules.
- ▶ Just a few months ago only Claude could sweep a codebase decently.
- ▶ Now both Claude and OpenAI can.
- ▶ And now they can prove like heck.
- ▶ 10x what I can do.
- ▶ Claude wants \$25 per code review now, probably because they can and it's expensive if they're using pure LLMs.

Software Engineering in AI Paired Proving- Problems

- ▶ What is the limiting factor now? **Code Review!**
- ▶ But really **Spec Review** when you prove.
- ▶ What can we do to simplify code review?
 - ▶ 1. Modularity - traits/impls in Verus.
 - ▶ 2. TOC - This organizes the boilerplate to make reading easier.
 - ▶ 3. Proof - I really just read the specs. If they're right the code is right!
 - ▶ 4. Tests - are hugely useful when you don't trust your coding team.
 - ▶ 5. Formatting - rust formatting has been adopted in Verus. It is very low density. I built a minimizing formatter to cut down on my working memory load while reviewing. F* is much tighter.

The AI Paired Programming Interfaces

- ▶ What you want is pretty obvious with experience: TEXT still rules.
- ▶ You want one window with your core interaction. The agent tells you what it is doing in it.
- ▶ You want one to watch your compile, tests and scripts.
- ▶ You want one window to have the LLM show you it's thinking and allow you to question and change it.
- ▶ I learned a ton from watching agents think.
- ▶ I suspect they are hiding most of their thinking now to prevent them from being used to train new AIs.
- ▶ And you want it not to delete your previous conversation on token compression.

Mythos Audit Rust/Firefox

#	Metric	Rust (Compiler/Std)	Firefox (Entire Codebase)
1	Codebase Size	~4,500 KLOC	~40,000 KLOC
2	Mythos Bugs Found	189	271
3	Vulnerability Density	0.0420 bugs/KLOC	0.0068 bugs/KLOC
4	Ratio	6.18	1

The Internet Apocalypse

- ▶ I asked Claude Code to check its switches on startup for me, at the End of Jan 2026.
- ▶ It immediately started disassembling itself.
- ▶ Scared the heck out of me.
- ▶ I told it to read it's docs.
- ▶ I have 100% belief that these LLMs are moving faster than computer security.
- ▶ It's verify or be hacked at this point.

The Internet Apocalypse

- ▶ Where should we focus our limited, but quickly growing, proving power?
- ▶ In theory 70% of intrusions are solved by type safe languages.
- ▶ I suspect that this just ignores way too much of the human aspect.
- ▶ I have never liked a single estimation paper's methods.
- ▶ One way to help us focus our new-found superpowers is to survey KLOC type-unsafe x Interfaces x Running Services x \$ Value.
- ▶ Which Amazon and Microsoft can probably measure.

Future Work

- ▶ GRASE (Medical Acronym) -
 - ▶ Generally Recognized as Safe and Effective
- ▶ **GRASE-**
 - ▶ **Git Recording Agentic Software Engineering**
- ▶ The GRASE project's goal is to make statistical process software optimization for agentic coding:
 - ▶ **Generally Regarded As Simple and Effective.**
- ▶ A Watt's Humphrey style PSP data collection done essentially effortlessly by the agents in the background.

Review of the Talk

- ▶ Ah let's just get to the questions!
- ▶ I have a ton of them for you! and I hope you have many for me.

Extra Slides

AI Paired Programming Size: Quick Review

- ▶ APAS-AI Lines of Code
 - ▶ Source: 31,751
 - ▶ Tests: 55,223
 - ▶ Benches: 13,890
 - ▶ Total code: 100,864
- ▶ Rusticate Lines of Code
 - ▶ Source: 29,582
 - ▶ Tests: 12,890
 - ▶ Total code: 42,472

AI Paired Proving Size: Quick Review

- ▶ APAS-VERUS Lines of Code
 - ▶ Source: 229,319
 - ▶ Tests (RTT): 58,503
 - ▶ PTTs (rust_verify_test): 15,938
 - ▶ Benches: 8,027
 - ▶ Total code: 311,787
- ▶ Veracity Lines of Code
 - ▶ Source: 65,189
 - ▶ Tests: 2,851
 - ▶ Total code: 68,040

Blood on the Road - Open Source =

- ▶ Open Source is now Open Copy.
- ▶ with reports of 1/2 of AI time
- ▶ and with the US Government totally undermining copyright law.

Blood on the Road - Open Repositories =

- ▶ Open repositories are now weakly locked doors.
- ▶ Agents can and have been used to inject invisible UNICODE attacks.
- ▶ Agents can now pass the Turing test, and more easily the programmer Turing test.
- ▶ They'll submit code pretending to be human.
- ▶ They'll steal credentials and submit code.
- ▶ LiteLLM is the most atrocious case lately.

Blood on the Road - Key Services

- ▶ Authentication is going to be attacked brutally.
- ▶ Network services are going to be hacked left and right.
- ▶ Configurations are going to be heavily attacked.
 - ▶ Using text files for these is already error prone.

Blood on the Road - Open Binaries =

- ▶ Open binaries are next.
- ▶ They are going to be disassembled and reconstructed.
- ▶ They are going to be rewritten in real time.
- ▶ Key solutions to this are:
- ▶ Binary packages from all repositories.
- ▶ Cryptographically signed.
- ▶ Installation protected in the OS.
- ▶ And protected from reading by the OS.

Questions - F* and Pulse

- ▶ Are F* modules really in use? If not, why?
- ▶ Got Functors?
- ▶ Are F* modules working with proof smoothly? Took me about three days to get a proving model (Treap) with modules.
- ▶ Does F* have a typed library search yet?
- ▶ Are typeclasses being heavily used?
- ▶ How fast is validation now?
- ▶ Nik you have any more statistics on AutoCLRS?

Questions

- ▶ Did the PL community overly complicate things with higher-order types?
- ▶ What can you folks say about MSFT adoption of proven code?
- ▶ What are the biggest F* verified systems?
- ▶ What are the biggest Pulse verified systems?
- ▶ Any quantitatatives? Person days? AI days?
- ▶ Should I quit Verus and go learn Pulse?

Questions

- ▶ Are you folks still using Make? That's how I got started on this.
- ▶ Many programmers and worse, managers, complain about the programmer time to switch languages. Are you finding that you can switch languages with your AIs explaining things to you?
- ▶ Can we prove Rust's (awful) std with Verus?
- ▶ Can we prove a compiler with Verus and EPR?
- ▶ Can we prove a TIL (Tarditi, Morrisett, Harper) like compiler?
- ▶ Can we verify a comp cert like compiler?
- ▶ Can we prove an OS in Verus?

Questions - AI LLM Agents

- ▶ How fast is this moving?
- ▶ How much context do we need?
- ▶ How much speed do we have? T/S down and up?
- ▶ How much speed do we need?
- ▶ How much speed will we get?
- ▶ When it's 10 times faster how will we use it?
- ▶ When will this really cost the user what it costs the provider?
- ▶ When will it get even better intermediate term memory?

Cloud Security Alliance

- ▶ The cloud security alliance has some sober recommendations:
- ▶ The “AI Vulnerability Storm”: Building a “Mythosready” Security Program
- ▶ <https://labs.cloudsecurityalliance.org/wp-content/uploads/2026/04/mythosreadyv91.pdf>
- ▶ Use LLM-based vulnerability discovery and remediation capabilities.
- ▶ Update risk metrics.
- ▶ We cannot outwork machine-speed threats. Re-prioritize, automate, and prepare for burnout.